

MFE R Programming Workshop

Week 6

Brett Dunn and Mahyar Kargar

Fall 2017

Introduction

Questions

Any questions before we start?

Overview

- ▶ `%>%`
- ▶ `tidyr`
- ▶ `dplyr`

$\%>\%$

The Pipe Operator %>%

- ▶ The `magrittr` package provides a pipe operator.
- ▶ See `vignette("magrittr")`.
- ▶ Basic piping:
 - ▶ `x %>% f` is equivalent to `f(x)`
 - ▶ `x %>% f(y)` is equivalent to `f(x, y)`
 - ▶ `x %>% f %>% g %>% h` is equivalent to `h(g(f(x)))`
- ▶ The argument placeholder:
 - ▶ `x %>% f(y, .)` is equivalent to `f(y, x)`
 - ▶ `x %>% f(y, z = .)` is equivalent to `f(y, z = x)`

Expose the variables with %\$%

- ▶ The %\$% allows variable names (e.g. column names) to be used in a function.

```
library(magrittr)
iris %>%
  subset(Sepal.Length > mean(Sepal.Length)) %$%
  cor(Sepal.Length, Sepal.Width)
```

```
## [1] 0.3361992
```

Compound assignment pipe operations with %<>%

- ▶ There is also a pipe operator which can be used as shorthand notation in situations where the left-hand side is being “overwritten”:

```
iris$Sepal.Length <-  
  iris$Sepal.Length %>%  
  sqrt()
```

Use the %<>% operator to avoid the repetition:

```
iris$Sepal.Length %<>% sqrt
```

- ▶ This operator works exactly like %>%, except the pipeline assigns the result rather than returning it.

tidyr

Hadley Wickham

- ▶ [Hadley Wickham](#) is practically famous in the R world
- ▶ He's developed a very large number of useful packages, e.g. `ggplot2` and `lubridate`.
- ▶ Today we will look at `dplyr` and `tidyr`.
- ▶ Tidy data is data that's easy to work with: it's easy to munge (with `dplyr`), visualise (with `ggplot2` or `ggvis`) and model (with R's hundreds of modelling packages).
- ▶ The two most important properties of tidy data are:
 - ▶ Each column is a variable.
 - ▶ Each row is an observation.
- ▶ Check [R for Data Science](#) book.

Sample data

- ▶ A common problem is a dataset where some of the column names are not names of variables, but values of a variable.
- ▶ Take table4a: the column names 1999 and 2000 represent values of the year variable, and each row represents two observations, not one.
- ▶ tidyr is a member of the core tidyverse.

```
library(tidyverse)
table4a
```

```
## # A tibble: 3 × 3
##       country `1999` `2000`
## *      <chr>   <int>   <int>
## 1 Afghanistan     745     2666
## 2      Brazil  37737    80488
## 3        China 212258   213766
```

Bring columns together with `gather()`

- ▶ To tidy a dataset like this, we need to **gather** those columns into a new pair of variables. To describe that operation we need three parameters:
 - ▶ The set of columns that represent values, not variables. In this example, those are the columns 1999 and 2000.
 - ▶ The name of the variable whose values form the column names. I call that the `key`, and here it is `year`.
 - ▶ The name of the variable whose values are spread over the cells. I call that `value`, and here it's the number of cases.

Bring columns together with `gather()`

- In the final result, the gathered columns are dropped, and we get new key and value columns.

```
table4a %>%  
  gather(`1999`, `2000`, key = "year", value = "cases")
```

```
## # A tibble: 6 × 3  
##       country year  cases  
##       <chr> <chr> <int>  
## 1 Afghanistan 1999    745  
## 2      Brazil 1999  37737  
## 3      China 1999 212258  
## 4 Afghanistan 2000   2666  
## 5      Brazil 2000  80488  
## 6      China 2000 213766
```

Split a column with spread()

- ▶ Spreading is the opposite of gathering. You use it when an observation is scattered across multiple rows.
- ▶ For example, take table2: an observation is a country in a year, but each observation is spread across two rows.

```
table2
```

```
## # A tibble: 12 × 4
##       country year      type      count
##       <chr> <int>    <chr>    <int>
## 1 Afghanistan 1999    cases      745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000    cases      2666
## 4 Afghanistan 2000 population 20595360
## 5      Brazil 1999    cases      37737
## 6      Brazil 1999 population 172006362
## 7      Brazil 2000    cases      80488
## 8      Brazil 2000 population 174504898
## 9      Brazil 2000    cases      812252
```

spreading

- ▶ To tidy this up, we first analyse the representation in similar way to `gather()`. This time, however, we only need two parameters:
 - ▶ The column that contains variable names, the `key` column. Here, it's `type`.
 - ▶ The column that contains values forms multiple variables, the `value` column. Here it's `count`.

spreading

- Once we've figured that out, we can use `spread()`, as shown below

```
spread(table2, key = type, value = count)
```

```
## # A tibble: 6 × 4
##   country year cases population
## *   <chr> <int> <int>      <int>
## 1 Afghanistan 1999    745  19987071
## 2 Afghanistan 2000   2666  20595360
## 3      Brazil 1999  37737  172006362
## 4      Brazil 2000  80488  174504898
## 5       China 1999 212258 1272915272
## 6       China 2000 213766 1280428583
```


spread() and gather() are complements

```
df <- data.frame(x = c("a", "b"), y = c(3, 4),  
                 z = c(5, 6))
```

```
df
```

```
##    x y z  
## 1 a 3 5  
## 2 b 4 6
```

```
df %>% spread(x, y) %>% gather(x, y, a:b, na.rm = TRUE)
```

```
##    z x y  
## 1 5 a 3  
## 4 6 b 4
```

There's much more

- ▶ As usual, read the [vignette](#) on the CRAN page
- ▶ Also check [Chapter 12](#) of R for Data Science book.

dplyr

Overview of dplyr

- ▶ dplyr provides a grammar of data manipulation.
 - ▶ A simple way to interact with data.
- ▶ We learn about:
 - ▶ tibble structure `tbl`
 - ▶ The pipe operator `%>%`
 - ▶ Using dplyr with databases
- ▶ The [dplyr introduction vignette](#) is a good resource.

dplyr and data.table

- ▶ See this [post](#).
- ▶ Here are my thoughts:
 - ▶ For data less than 1 million rows, it is reported that there is not a significant speed difference between the two.
 - ▶ For large data that can fit in memory, use `data.table`.
 - ▶ For data that cannot fit in memory, you could use `dplyr` with a database backend.
- ▶ `dtplyr` is a package to use `dplyr` with `data.table`.
 - ▶ It is slower than just using `data.table`.

Data: nycflights13

- ▶ To explore the basic data manipulation verbs of dplyr, we'll start with the built in 'nycflights13' data frame
- ▶ This dataset contains all flights that departed from New York City in 2013

```
library(dplyr)
library(nycflights13)
```

```
head(flights,4)
```

```
## # A tibble: 4 × 19
##   year month   day dep_time sched_dep_time dep_delay
##   <int> <int> <int>   <int>         <int>         <dbl>
## 1  2013     1     1     517             515           2
## 2  2013     1     1     533             529           4
## 3  2013     1     1     542             540           2
## 4  2013     1     1     544             545          -1
## # ... with 13 more variables: arr_time <int>,
```

tbls (Tibbles)

- ▶ A tbl will only display the data that will fit in your console.
-glimpse() is another nice way to look at the data

```
flights <- tbl_df(flights)
flights
```

```
## # A tibble: 336,776 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay
##   <int> <int> <int>   <int>         <int>         <dbl>
## 1  2013     1     1     517           515           2
## 2  2013     1     1     533           529           4
## 3  2013     1     1     542           540           2
## 4  2013     1     1     544           545          -1
## 5  2013     1     1     554           600          -6
## 6  2013     1     1     554           558          -4
## 7  2013     1     1     555           600          -5
## 8  2013     1     1     557           600          -3
## 9  2013     1     1     557           600          -3
```

Single Table Verbs

- ▶ `dplyr` aims to provide a function for each basic verb of data manipulation:
- ▶ `select()` (and `rename()`)
 - ▶ returns a subset of the columns
- ▶ `filter()` (and `slice()`)
 - ▶ returns a subset of the rows
- ▶ `arrange()` - reorders rows
 - ▶ reorders the rows according to single or multiple variables
- ▶ `distinct()`
- ▶ `mutate()` (and `transmute()`)
 - ▶ builds adds new columns from the data
- ▶ `summarise()` - calculates summary statistics
 - ▶ which reduces each group to a single row by calculating aggregate measures
- ▶ `sample_n()` and `sample_frac()`

Tidy Data

- ▶ `dplyr` works best when variables are in columns and observations are in rows.
- ▶ You can use `tidyr` to help you create a tidy dataset.

Select Columns by Name with `select()`

- ▶ `select()` allows you to rapidly zoom in on a useful subset using operations that usually only work on numeric variable positions:

```
# Select columns by name  
select(flights, year, month, day)
```

```
## # A tibble: 336,776 × 3
```

```
##   year month   day
```

```
##   <int> <int> <int>
```

```
## 1  2013     1     1
```

```
## 2  2013     1     1
```

```
## 3  2013     1     1
```

```
## 4  2013     1     1
```

```
## 5  2013     1     1
```

```
## 6  2013     1     1
```

```
## 7  2013     1     1
```

```
## 8  2013     1     1
```

Select a Range of Columns with :

```
# Select all columns between year and day (inclusive)  
select(flights, year:day)
```

```
## # A tibble: 336,776 × 3  
##   year month   day  
##   <int> <int> <int>  
## 1  2013     1     1  
## 2  2013     1     1  
## 3  2013     1     1  
## 4  2013     1     1  
## 5  2013     1     1  
## 6  2013     1     1  
## 7  2013     1     1  
## 8  2013     1     1  
## 9  2013     1     1  
## 10 2013     1     1  
## # ... with 336,766 more rows
```

An Example of `-(col1:col2)`

```
# Select all columns except those from year to day (inclus  
select(flights, -(year:day))
```

```
## # A tibble: 336,776 × 16
```

```
##   dep_time sched_dep_time dep_delay arr_time
```

```
##   <int>         <int>         <dbl>    <int>
```

```
## 1      517           515           2      830
```

```
## 2      533           529           4      850
```

```
## 3      542           540           2      923
```

```
## 4      544           545          -1     1004
```

```
## 5      554           600          -6      812
```

```
## 6      554           558          -4      740
```

```
## 7      555           600          -5      913
```

```
## 8      557           600          -3      709
```

```
## 9      557           600          -3      838
```

```
## 10     558           600          -2      753
```

```
## # ... with 336,766 more rows, and 12 more variables:
```

```
## #   sched arr time <int>   arr delay <dbl>   carrier <chr>
```

select Helper Functions

- ▶ `dplyr` comes with a set of helper functions that can help you select groups of variables inside a `select()` call:
- ▶ `starts_with("X")`: every name that starts with "X",
- ▶ `ends_with("X")`: every name that ends with "X",
- ▶ `contains("X")`: every name that contains "X",
- ▶ `matches("X")`: every name that matches "X", where "X" can be a regular expression,
- ▶ `num_range("x", 1:5)`: the variables named `x01`, `x02`, `x03`, `x04` and `x05`,
- ▶ `one_of(x)`: every name that appears in `x`, which should be a character vector.

Add New Columns with mutate()

```
mutate(flights,  
       gain = arr_delay - dep_delay,  
       speed = distance / air_time * 60)
```

```
## # A tibble: 336,776 × 21
```

```
##   year month   day dep_time sched_dep_time dep_delay  
##   <int> <int> <int>   <int>         <int>         <dbl>  
## 1  2013     1     1     517             515           2  
## 2  2013     1     1     533             529           4  
## 3  2013     1     1     542             540           2  
## 4  2013     1     1     544             545          -1  
## 5  2013     1     1     554             600          -6  
## 6  2013     1     1     554             558          -4  
## 7  2013     1     1     555             600          -5  
## 8  2013     1     1     557             600          -3  
## 9  2013     1     1     557             600          -3  
## 10 2013     1     1     558             600          -2
```

```
## # ... with 336,766 more rows, and 15 more variables
```

If you only want to keep the new variables, use `transmute()`

```
transmute(flights,  
          gain = arr_delay - dep_delay,  
          gain_per_hour = gain / (air_time / 60)  
)
```

```
## # A tibble: 336,776 × 2  
##       gain gain_per_hour  
##   <dbl>      <dbl>  
## 1      9      2.378855  
## 2     16      4.229075  
## 3     31     11.625000  
## 4    -17     -5.573770  
## 5    -19     -9.827586  
## 6     16      6.400000  
## 7     24      9.113924  
## 8    -11    -12.452830  
## 9      5      2.142857
```

Filter rows with `filter()`

- ▶ `filter()` allows you to select a subset of rows in a data frame.
- ▶ The first argument is the name of the data frame.
- ▶ The second and subsequent arguments are the expressions that filter the data frame
- ▶ Select all flights on January 1st with:

```
filter(flights, month == 1, day == 1)
```

```
## # A tibble: 842 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay
##   <int> <int> <int>   <int>         <int>         <dbl>
## 1  2013     1     1     517             515           2
## 2  2013     1     1     533             529           4
## 3  2013     1     1     542             540           2
## 4  2013     1     1     544             545          -1
## 5  2013     1     1     554             600          -6
## 6  2013     1     1     554             558          -4
## 7  2013     1     1     555             600          -5
```


Select rows by position

- To select rows by position, use `slice()`

```
slice(flights, 1:10)
```

```
## # A tibble: 10 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay  
##   <int> <int> <int>   <int>         <int>         <dbl>  
## 1  2013     1     1     517             515           2  
## 2  2013     1     1     533             529           4  
## 3  2013     1     1     542             540           2  
## 4  2013     1     1     544             545          -1  
## 5  2013     1     1     554             600          -6  
## 6  2013     1     1     554             558          -4  
## 7  2013     1     1     555             600          -5  
## 8  2013     1     1     557             600          -3  
## 9  2013     1     1     557             600          -3  
## 10 2013     1     1     558             600          -2  
## # ... with 13 more variables: arr_time <int>,
```

Arrange rows with `arrange()`

- ▶ `arrange()` works similarly to `filter()` except that instead of filtering or selecting rows, it reorders them.

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay
##   <int> <int> <int>   <int>         <int>         <dbl>
## 1  2013     1     1     517           515           2
## 2  2013     1     1     533           529           4
## 3  2013     1     1     542           540           2
## 4  2013     1     1     544           545          -1
## 5  2013     1     1     554           600          -6
## 6  2013     1     1     554           558          -4
## 7  2013     1     1     555           600          -5
## 8  2013     1     1     557           600          -3
## 9  2013     1     1     557           600          -3
## 10 2013     1     1     558           600          -2
```

Use desc() to order a column in descending order

```
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>
```

```
## 1  2013     1     9     641           900         1301
```

```
## 2  2013     6    15    1432          1935         1137
```

```
## 3  2013     1    10    1121          1635         1126
```

```
## 4  2013     9    20    1139          1845         1014
```

```
## 5  2013     7    22     845          1600         1005
```

```
## 6  2013     4    10    1100          1900          960
```

```
## 7  2013     3    17    2321           810          911
```

```
## 8  2013     7    22    2257           759          898
```

```
## 9  2013    12     5     756          1700          896
```

```
## 10 2013     5     3    1133          2055          878
```

```
## # ... with 336,766 more rows, and 13 more variables:
```

```
## #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>
```

You can rename variables with `rename()`

```
rename(flights, tail_num = tailnum)
```

```
## # A tibble: 336,776 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>
```

```
## 1  2013     1     1     517           515           2
```

```
## 2  2013     1     1     533           529           4
```

```
## 3  2013     1     1     542           540           2
```

```
## 4  2013     1     1     544           545          -1
```

```
## 5  2013     1     1     554           600          -6
```

```
## 6  2013     1     1     554           558          -4
```

```
## 7  2013     1     1     555           600          -5
```

```
## 8  2013     1     1     557           600          -3
```

```
## 9  2013     1     1     557           600          -3
```

```
## 10 2013     1     1     558           600          -2
```

```
## # ... with 336,766 more rows, and 13 more variables:
```

```
## #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>
```

```
## #   carrier <chr>, flight <int>, tail_num <chr>
```

Extract distinct (unique) rows

- ▶ A common use of `select()` is to find the values of a set of variables.
- ▶ This is particularly useful in conjunction with the `distinct()` verb

```
distinct(select(flights, tailnum))
```

```
## # A tibble: 4,044 × 1
```

```
##   tailnum
```

```
##   <chr>
```

```
## 1  N14228
```

```
## 2  N24211
```

```
## 3  N619AA
```

```
## 4  N804JB
```

```
## 5  N668DN
```

```
## 6  N39463
```

```
## 7  N516JB
```

```
## 8  N829AS
```

Summarise values with summarise()

- ▶ The last verb is `summarise()`. It collapses a data frame to a single row.
- ▶ You can use any function you like in `summarise()` so long as the function can take a vector of data and return a single number.

```
summarise(flights,  
          delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 × 1  
##       delay  
##       <dbl>  
## 1 12.63907
```

dplyr aggregate functions

- ▶ dplyr provides several helpful aggregate functions of its own, in addition to the ones that are already defined in R. These include:
 - ▶ `first(x)` - The first element of vector `x`.
 - ▶ `last(x)` - The last element of vector `x`.
 - ▶ `nth(x, n)` - The `n`th element of vector `x`.
 - ▶ `n()` - The number of rows in the data.frame or group of observations that `summarise()` describes.
 - ▶ `n_distinct(x)` - The number of unique values in vector `x`.

Chaining

- ▶ The dplyr API is functional — function calls don't have side-effects.
- ▶ You must always save their results. **UGLY**
- ▶ To get around this problem, dplyr provides the `%>%` operator
- ▶ `x %>% f(y)` turns into `f(x, y)`

```
flights %>%  
group_by(year, month, day) %>%  
select(arr_delay, dep_delay) %>%  
summarise(arr = mean(arr_delay, na.rm = TRUE),  
dep = mean(dep_delay, na.rm = TRUE)) %>%  
filter(arr > 30 | dep > 30)
```

```
## Adding missing grouping variables: `year`, `month`, `day`
```

```
## Source: local data frame [49 x 5]
```

```
## Groups: year, month [11]
```

```
##
```


Commonalities

- ▶ The syntax and function of all these verbs are very similar:
- ▶ The first argument is a data frame.
- ▶ The subsequent arguments describe what to do with the data frame.
- ▶ The result is a new data frame
- ▶ Together these properties make it easy to chain together multiple simple steps to achieve a complex result.

Grouped operations

- ▶ These verbs are useful on their own, but they become really powerful when you apply them to groups of observations
- ▶ In dplyr, you do this by with the `group_by()` function.
- ▶ It breaks down a dataset into specified groups of rows.

Grouped operations (cont.)

- ▶ Grouping affects the verbs as follows:
- ▶ `grouped select()` is the same as `ungrouped select()`, except that grouping variables are always retained.
- ▶ `grouped arrange()` orders first by the grouping variables
- ▶ `mutate()` and `filter()` are most useful in conjunction with window functions (like `rank()`, or `min(x) = x`). They are described in detail in `vignette("window-functions")`.
- ▶ `sample_n()` and `sample_frac()` sample the specified number/fraction of rows in each group.
- ▶ `slice()` extracts rows within each group.
- ▶ `summarise()` is powerful and easy to understand, as described in more detail below.

group_by Example

- For example, we could use these to find the number of planes and the number of flights that go to each possible destination:

```
flights %>%  
  group_by(dest) %>%  
  summarise(planes = n_distinct(tailnum),  
            flights = n())
```

```
## # A tibble: 105 × 3  
##   dest planes flights  
##   <chr>   <int>   <int>  
## 1 ABQ     108     254  
## 2 ACK      58     265  
## 3 ALB     172     439  
## 4 ANC       6       8  
## 5 ATL    1180    17215  
## 6 AUS     993    2439  
## 7 AVL     159     275
```

Multiple table verbs

- ▶ dplyr implements the four most useful SQL joins:
- ▶ `inner_join(x, y)`: matching $x + y$
- ▶ `left_join(x, y)`: all $x +$ matching y
- ▶ `semi_join(x, y)`: all x with match in y
- ▶ `anti_join(x, y)`: all x without match in y
- ▶ And provides methods for:
- ▶ `intersect(x, y)`: all rows in both x and y
- ▶ `union(x, y)`: rows in either x or y
- ▶ `setdiff(x, y)`: rows in x , but not y

Joins from dplyr Map to SQL

- ▶ `inner_join(x, y)`
 - ▶ `SELECT * FROM x JOIN y ON x.a = y.a`
- ▶ `left_join(x, y)`
 - ▶ `SELECT * FROM x LEFT JOIN y ON x.a = y.a`
- ▶ `right_join(x, y)`
 - ▶ `SELECT * FROM x RIGHT JOIN y ON x.a = y.a`
- ▶ `full_join(x, y)`
 - ▶ `SELECT * FROM x FULL JOIN y ON x.a = y.a`
- ▶ `semi_join(x, y)`
 - ▶ `SELECT * FROM x WHERE EXISTS (SELECT 1 FROM y WHERE x.a = y.a)`
- ▶ `anti_join(x, y)`
 - ▶ `SELECT * FROM x WHERE NOT EXISTS (SELECT 1 FROM y WHERE x.a = y.a)`

Lab 3

- ▶ Let's redo [Lab 3](#) with dplyr.